

Assignments

This page will contain all the assignments you submit for the class.

Instructions for all assignments

I want you to submit your assignment as a PDF, so I can keep a record of what the code looked like that day. I also want you to include your answers on your personal GitHub website. This will be good practice for editing your website and it will help you produce something you can keep after the class is over.

1. Download the Assignment1.Rmd file from Canvas. You can use this as a template for writing your answers. It's the same as what you can see on my website in the Assignments tab. Once we're done with this I'll edit the text on the website to include the solutions.
2. On RStudio, open a new R script in RStudio (File > New File > R Script). This is where you can test out your R code. You'll write your R commands and draw plots here.
3. Once you have finalized your code, copy and paste your results into this template (Assignment 1.Rmd). For example, if you produced a plot as the solution to one of the problems, you can copy and paste the R code in R markdown by using the ```{r} ``` command. Answer the questions in full sentences and Save.
4. Produce a PDF file with your answers. To do this, knit to PDF (use Knit button at the top of RStudio), locate the PDF file in your docs folder (it's in the same folder as the Rproj), and submit that on on Canvas in Assignment 1.
5. Build Website, go to GitHub desktop, commit and push. Now your solutions should be on your website as well.

Assignment 1

Collaborators:

Problem 1

Install the datasets package on the console below using `install.packages("datasets")`. Now load the library.

```
#install.packages("datasets")  
library(datasets)
```

```
USArrests
```

##		Murder	Assault	UrbanPop	Rape
##	Alabama	13.2	236	58	21.2
##	Alaska	10.0	263	48	44.5
##	Arizona	8.1	294	80	31.0
##	Arkansas	8.8	190	50	19.5
##	California	9.0	276	91	40.6
##	Colorado	7.9	204	78	38.7
##	Connecticut	3.3	110	77	11.1
##	Delaware	5.9	238	72	15.8
##	Florida	15.4	335	80	31.9
##	Georgia	17.4	211	60	25.8
##	Hawaii	5.3	46	83	20.2
##	Idaho	2.6	120	54	14.2
##	Illinois	10.4	249	83	24.0
##	Indiana	7.2	113	65	21.0
##	Iowa	2.2	56	57	11.3
##	Kansas	6.0	115	66	18.0
##	Kentucky	9.7	109	52	16.3
##	Louisiana	15.4	249	66	22.2
##	Maine	2.1	83	51	7.8
##	Maryland	11.3	300	67	27.8
##	Massachusetts	4.4	149	85	16.3
##	Michigan	12.1	255	74	35.1
##	Minnesota	2.7	72	66	14.9
##	Mississippi	16.1	259	44	17.1
##	Missouri	9.0	178	70	28.2
##	Montana	6.0	109	53	16.4
##	Nebraska	4.3	102	62	16.5
##	Nevada	12.2	252	81	46.0
##	New Hampshire	2.1	57	56	9.5
##	New Jersey	7.4	159	89	18.8
##	New Mexico	11.4	285	70	32.1
##	New York	11.1	254	86	26.1
##	North Carolina	13.0	337	45	16.1
##	North Dakota	0.8	45	44	7.3
##	Ohio	7.3	120	75	21.4
##	Oklahoma	6.6	151	68	20.0
##	Oregon	4.9	159	67	29.3
##	Pennsylvania	6.3	106	72	14.9
##	Rhode Island	3.4	174	87	8.3
##	South Carolina	14.4	279	48	22.5
##	South Dakota	3.8	86	45	12.8
##	Tennessee	13.2	188	59	26.9
##	Texas	12.7	201	80	25.5
##	Utah	3.2	120	80	22.9
##	Vermont	2.2	48	32	11.2
##	Virginia	8.5	156	63	20.7
##	Washington	4.0	145	73	26.2
##	West Virginia	5.7	81	39	9.3
##	Wisconsin	2.6	53	66	10.8
##	Wyoming	6.8	161	60	15.6

Load the USArrests dataset and rename it `dat`. Note that this dataset comes with R, in the package `datasets`, so there's no need to load data from your computer. Why is it useful to rename the dataset?

It is useful to renamed USArrests to dat.us because it is easier to write and it is good practice to rewrite data for yourself so you can create your own data which can be replicated by another person if they use the original data.

```
dat.us <- USArrests  
  
head(dat.us)
```

```
##           Murder Assault UrbanPop Rape  
## Alabama      13.2      236        58 21.2  
## Alaska       10.0      263        48 44.5  
## Arizona       8.1      294        80 31.0  
## Arkansas      8.8      190        50 19.5  
## California    9.0      276        91 40.6  
## Colorado      7.9      204        78 38.7
```

Problem 2

Use this command to make the state names into a new variable called State.

```
dat.us$state <- tolower(rownames(USArrests))
```

This dataset has the state names as row names, so we just want to make them into a new variable. We also make them all lower case, because that will help us draw a map later - the map function requires the states to be lower case.

List the variables contained in the dataset USArrests.

```
names(dat.us)  
  
## [1] "Murder"  "Assault" "UrbanPop" "Rape"     "state"
```

Answer: The four variables are Murder, Assault, UrbanPop, Rape.

Problem 3

What type of variable (from the DVB chapter) is Murder?

Answer: Murder is a quantitative variable.

What R Type of variable is it?

Answer: Murder is numeric.

Problem 4


What information is contained in this dataset, in general? What do the numbers mean?

Answer: The dataset contains the data of murder, assault, rape and urbanpop from all 50 US states. The numbers represent the frequency of arrests for one of the four variables in a state during the time frame that the data was collected.

Problem 5

Draw a histogram of `Murder` with proper labels and title.

```
hist(dat.us$Murder, xlab= "Murders", ylab="Frequency", main= "Murder in the US", xlim=(c(0, 20)), ylim=
```



Problem 6

Please summarize `Murder` quantitatively. What are its mean and median? What is the difference between mean and median? What is a quartile, and why do you think R gives you the 1st Qu. and 3rd Qu.?

```
summary(dat.us$Murder)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.800   4.075   7.250   7.788  11.250   17.400
```

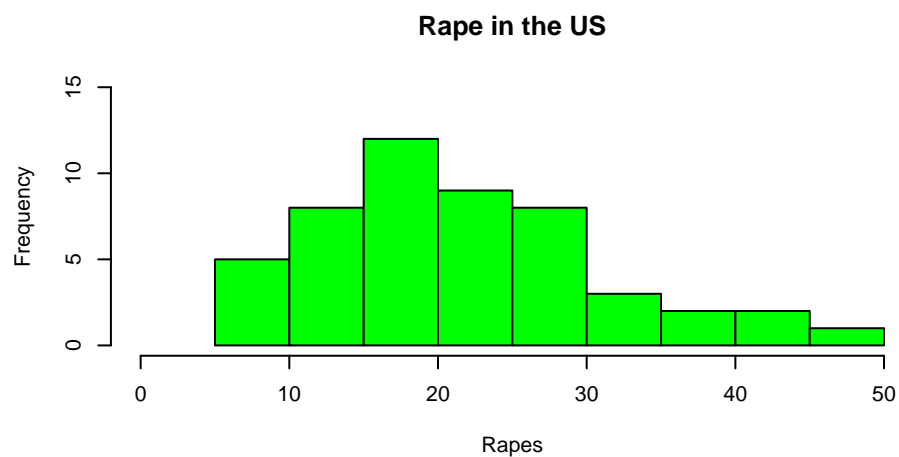
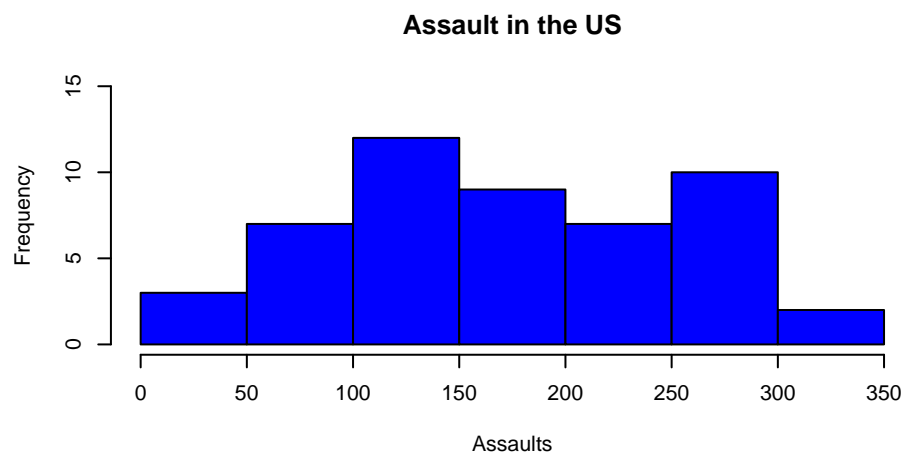

Answer: The mean for `Murder` is 7.788 and the median is 7.250. The mean is the amount that each subject would have if all of the values were added together and evenly distributed. If all 50 states had the same frequency of arrests for murders then it would be 7.788. The median is the middle value where exactly 50% of the values fall either above or below it. In the US, 50% of states have an arrest for murder frequency above 7.250 and the other 50% is below that. The median is highly robust because it is not greatly affected by outliers. The mean is the most common measure of central tendency but it is not robust because it will change based on the skewness of the distribution. A quartile indicates an interval that contains 25% or a quarter of the data. The first quartile for “Murder” is 4.075 which means that 25% of the “Murder” data falls

below 4.075 and the 3rd quartile is 11.250 which means that 25% of US states have a frequency of arrests for murder that is higher than 11.250. R gives you the 1st and 3rd quartile because those values are useful in determining the interquartile range (IQR). The IQR is the central half which means that 50% of the data falls within the 1st and 3rd quartile. In a box plot, values 1.5 IQRs above or below the tails are considered outliers.

Problem 7

Repeat the same steps you followed for `Murder`, for the variables `Assault` and `Rape`. Now plot all three histograms together. You can do this by using the command `par(mfrow=c(3,1))` and then plotting each of the three.

```
par(mfrow=c(3,1))  
  
hist(dat.us$Murder, xlab= "Murders", ylab="Frequency", main= "Murder in the US", xlim=(c(0, 20)), ylim=  
hist(dat.us$Assault, xlab= "Assaults", ylab="Frequency", main= "Assault in the US", xlim=(c(0, 350)), y  
hist(dat.us$Rape, xlab= "Rapes", ylab="Frequency", main= "Rape in the US", xlim=(c(0, 50)), ylim=(c(0, 1
```



What does the command `par` do, in your own words (you can look this up by asking R `?par`)?

Answer: Command `par` is used to set parameters. The `mfrow` input allows you to create an array to plot multiple graphs on one window. The command `par(mfrow=c(3,1))` allows three graphs to be plotted in three rows.

What can you learn from plotting the histograms together?

Answer: When the histograms are plotted together it is easier to compare the skewness and spread of each plot. You can see where each histogram has its peaks and outliers.

Problem 8

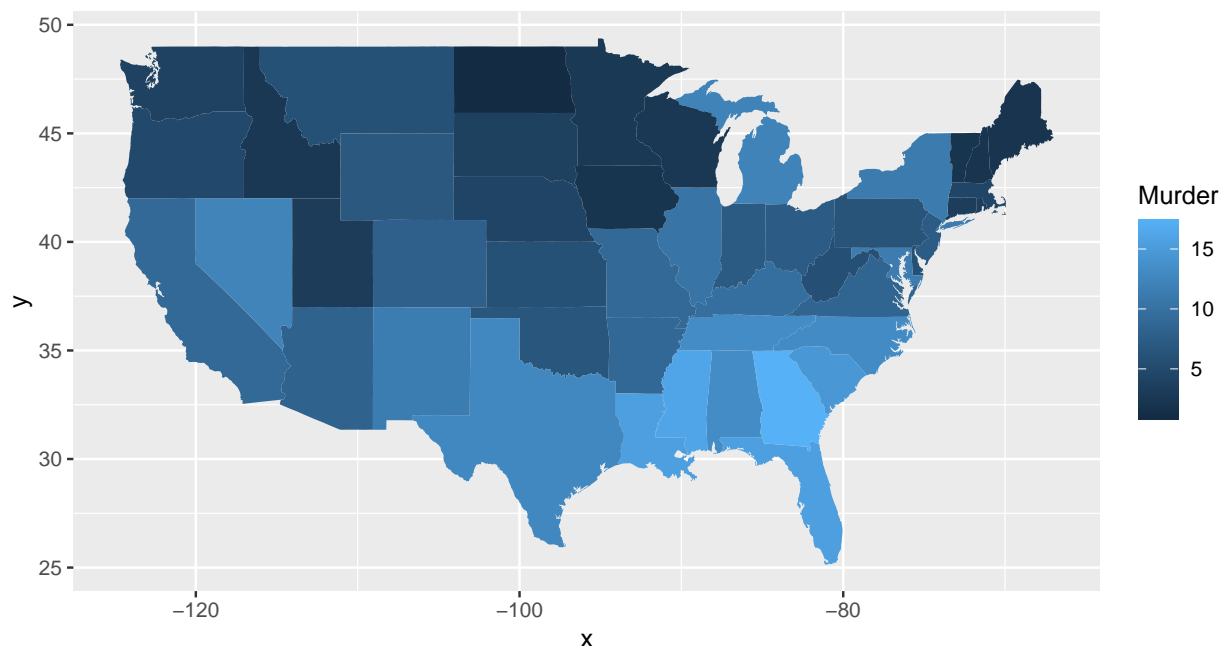
In the console below (not in text), type `install.packages("maps")` and press Enter, and then type `install.packages("ggplot2")` and press Enter. This will install the packages so you can load the libraries.

Run this code:

```
#install.packages("maps")
#install.packages("ggplot2")

library(maps)
library(ggplot2)

ggplot(dat.us, aes(map_id=state, fill=Murder)) +
  geom_map(map=map_data("state")) +
  expand_limits(x=map_data("state")$long, y=map_data("state")$lat)
```



What does this code do? Explain what each line is doing.

Answer: The lines `library('maps')` and `library('ggplot2')` are pulling from the packages that were installed. The line `ggplot(dat.us, aes(map_id=state, fill=Murder))` is creating a ggplot with the `USArrests` dataset. The plot is set with an aesthetic of a map with the US states. Each state is filled in with its respective murder arrest data. The line `geom_map(map=map_data("state"))` contains the map coordinates for each US state. The last line `'expand_limits(x=map_data("state")$long, y = map_data("state")$lat)'` ensures that the limits of the plot include a single value for all plots. The x and y axis of this plot contains the value of "state" from the map data. The x axis is longitude and the y axis is latitude. Together this code creates a

map of the US with each state filled in with its value for murder arrests. The darker blue indicates that the murder arrest frequency is 5 and below and the light blue indicates that it is 15 and above.

Assignment 2

Problem 1

```
dat <- read.csv(file = 'dat.nsduh.small.1.csv')
```

```
head(dat)
```

```
##   mjage cigage iralcage age2 sexatract speakengl irsex
## 1    14     50      14  16         1         1      1
## 2    11     14       5  13         2         1      2
## 3    12     35      12  15         2         1      2
## 4    16     18      18  14         1         1      1
## 5    14     16      14  16         4         1      1
## 6    12     16      18  15         4         1      2
```

```
names(dat)
```

```
## [1] "mjage"      "cigage"     "iralcage"   "age2"       "sexatract"  "speakengl"
## [7] "irsex"
```

```
dim(dat)
```

```
## [1] 171  7
```

What are the dimensions of the dataset?

Answer: The data has 7 columns for the 7 variables (mjage, cigage, iralcage, age2, sexatract, speakengl and irsex). There are 171 observations in the data.

Problem 2

Describe the variables in the dataset.

Answer: There are 7 variables in the dataset. `mjage` represents the age when the participants first used marijuana or hashish. `Cigage` represents the age when the participants first started smoking cigarettes everyday. `Iralcage` is the age when participants first tried alcohol. `Age2` is the final edited age of the participants. `Irsex` represents the gender of participants. `Sexattract` is the sexual attraction of the participants. `Speakeng` represents how well the participant speaks english.


What is this dataset about? Who collected the data, what kind of sample is it, and what was the purpose of generating the data?

Answer: This dataset is a small sample from the 2019 survey from the National Survey of Drug Use and Health about drug use in the 50 states and the District of Columbia in the United States. The survey was directed by the Substance Abuse and Mental Health Services Administration and it was conducted by RTI International. The survey is used to determine which populations and geographic areas have particular substance use problems so federal resources can be used effectively.

Problem 3: Age and gender

What is the age distribution of the sample like? Make sure you read the codebook to know what the variable values mean.

```
hist(dat$age2, xlab= "Final Edited Age", main= "Age Distribution of NSDUH 2019 Survey", col="purple")
```



*Answer:*The age distribution is negatively skewed because the bulk of the data is in the upper range with thin tails in the lower range. This shows that most of the respondents of the survey were 25 years old to greater than 65 years old. This age variable is categorical because the numbers represent categories of age for example 12 indicates that the respondent is 24 or 25 years old.


Do you think this age distribution representative of the US population? Why or why not?

*Answer:*I do not think this age distribution is representative of the US population because while the households chosen for the survey are randomly selected, only one person per household can respond to the survey and the interviews are completed online and not every person in the US has access to a computer or internet. There is also a money incentive to complete the survey. It is possible that more middle aged people responded to the survey.

Is the sample balanced in terms of gender? If not, are there more females or males?

Use this code to draw a stacked bar plot to view the relationship between sex and age. What can you conclude from this plot?

```
tab.agesex <- table(dat$irsex, dat$age2)
barplot(tab.agesex,
        main = "Stacked barchart",
        xlab = "Age category", ylab = "Frequency",
        legend.text = rownames(tab.agesex),
        beside = FALSE)
```



*Answer:*The sample is not fully balanced in terms of gender. There are more females with 52.28% of respondents being female and 47.72% being male.

Answer: From this plot, we can determine that the relationship between sex and age is fairly even. For the largest age group, the gender distribution appears to be almost equal. However, for the lower age groups there appears to be more female respondents and for the higher age distributions there appears to be more male respondents.

Problem 4: Substance use

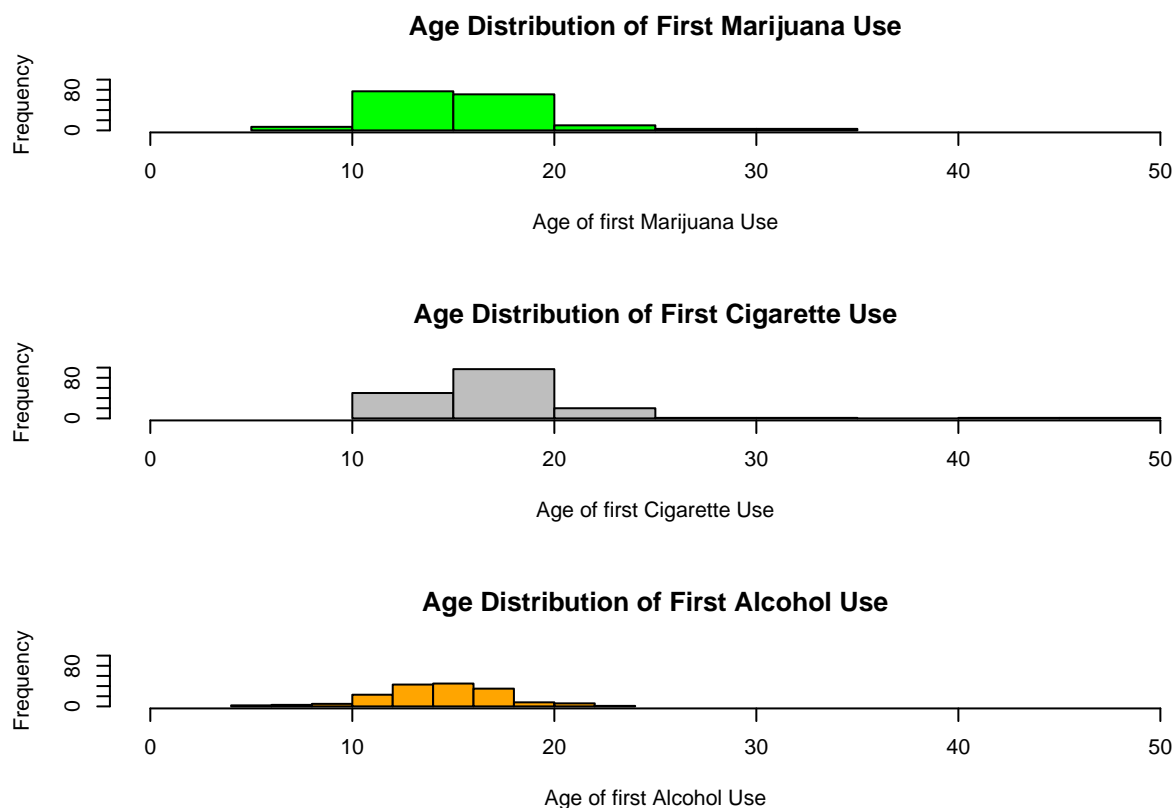
For which of the three substances included in the dataset (marijuana, alcohol, and cigarettes) do individuals tend to use the substance earlier?

```
par(mfrow=c(3,1))

hist(dat$mjage, col="green", xlab= "Age of first Marijuana Use", main= "Age Distribution of First Marijuana Use")

hist(dat$cigage, col="grey", xlab= "Age of first Cigarette Use", main= "Age Distribution of First Cigarette Use")

hist(dat$iralcage, col="orange", xlab= "Age of first Alcohol Use", main= "Age Distribution of First Alcohol Use")
```




Answer: Individuals tend to first start using marijuana and alcohol at a much younger age than cigarettes but individuals tend to first start using alcohol earliest.

Problem 5: Sexual attraction

What does the distribution of sexual attraction look like? Is this what you expected?


```
tab.sexatract <- table(dat$sexatract)
barplot(tab.sexatract, main= "Sexual Attraction Distribution", xlab= "Sexual Attraction", ylab= "Frequency")
```



Answer: Most participants are only attracted to the opposite sex. This is expected because respondents might not have felt comfortable answering this question unless they were only attracted to the same sex since there is stigma for different sexual attractions. The “99” category is significant because a large portion of respondents (23.87%) skipped this question which could indicate that people were not comfortable answering this question.

What is the distribution of sexual attraction by gender?


```
tab.sexatractsex <- table(dat$irsex, dat$sexatract)
barplot(tab.sexatractsex,
        main = "Stacked barchart",
        xlab = "Age category", ylab = "Frequency",
        legend.text = rownames(tab.sexatractsex),
        beside = FALSE)
```



Problem 6: English speaking

What does the distribution of English speaking look like in the sample? Is this what you might expect for a random sample of the US population?


```
tab.speakeng <- table(dat$speakengl)
barplot(tab.speakeng, xlab= "How Well English is Spoken", ylab= "Frequency", names.arg = c("Very Well",
```



Answer: Most of the participants speak English very well which is expected because for the most part the US is an English speaking country. It might have been hard for non-English speakers to engage with the survey if interpreters or different language options were not available.

Are there more English speaker females or males?

```
tab.speakengsex <- table(dat$irsex, dat$speakengl)
barplot(tab.speakengsex,
        main = "Stacked barchart",
        xlab = "Age category", ylab = "Frequency",
        legend.text = rownames(tab.speakengsex),
        beside = FALSE)
```



Answer: There appears to be an equal amount of English speaking females and males.

Exam 1

Load the data into an R data frame.

```
#read.csv(file = 'fatal-police-shootings-data.csv')  
dat <- read.csv(file = 'fatal-police-shootings-data.csv')  
  
head(dat)
```

```
##   id          name      date manner_of_death   armed age gender race
## 1 3      Tim Elliot 2015-01-02      shot      gun  53      M    A
## 2 4  Lewis Lee Lembke 2015-01-02      shot      gun  47      M    W
## 3 5 John Paul Quintero 2015-01-03 shot and Tasered  unarmed  23      M    H
## 4 8   Matthew Hoffman 2015-01-04      shot  toy weapon  32      M    W
## 5 9 Michael Rodriguez 2015-01-04      shot  nail gun  39      M    H
## 6 11 Kenneth Joe Brown 2015-01-04      shot      gun  18      M    W
##           city state signs_of_mental_illness threat_level      flee
## 1      Shelton   WA              True      attack Not fleeing
## 2       Aloha    OR              False     attack Not fleeing
## 3      Wichita   KS              False      other Not fleeing
## 4 San Francisco  CA              True      attack Not fleeing
## 5        Evans   CO              False     attack Not fleeing
## 6      Guthrie   OK              False     attack Not fleeing
## body_camera longitude latitude is_geocoding_exact
## 1      False  -123.122    47.247              True
## 2      False  -122.892    45.487              True
## 3      False   -97.281    37.695              True
## 4      False  -122.422    37.763              True
## 5      False  -104.692    40.384              True
## 6      False   -97.423    35.877              True
```

Problem 1

- Describe the dataset. This is the source: <https://github.com/washingtonpost/data-police-shootings> . Write two sentences (max.) about this.

This dataset contains the data of fatal police shootings of civilians in 2015. The data was collected from the Washington Post using local news reports, law enforcement websites, social media and independent databases.

- How many observations are there in the data frame?

```
dim(dat)
```

```
## [1] 6594  17
```

There are 6594 rows (observations) in the dataset with 17 columns.

- Look at the names of the variables in the data frame. Describe what “body_camera”, “flee”, and “armed” represent, according to the codebook. Again, only write one sentence (max) per variable.

```
names(dat)
```

```
## [1] "id"          "name"
## [3] "date"        "manner_of_death"
## [5] "armed"       "age"
## [7] "gender"      "race"
## [9] "city"        "state"
## [11] "signs_of_mental_illness" "threat_level"
## [13] "flee"        "body_camera"
## [15] "longitude"   "latitude"
## [17] "is_geocoding_exact"
```


The variable “body_camera indicates that the police officer was wearing a body camera and may have recorded some parts of the incident. The variable”flee” means that the victim was moving away from the officers. The variable “armed” means that victim had an instrument that the police believed could inflict harm.

- d. What are three weapons that you are surprised to find in the “armed” variable? Make a table of the values in “armed” to see the options.

```
table(dat$armed)
```

```
##
##
##          207          air conditioner
##          air pistol      Airsoft pistol
##          1              3
##          ax              barstool
##          24              1
##          baseball bat    baseball bat and bottle
##          20              1
## baseball bat and fireplace poker    baseball bat and knife
##          1              1
##          baton          BB gun
##          6              15
##          BB gun and vehicle    bean-bag gun
##          1              1
##          beer bottle          binoculars
##          3              1
##          blunt object          bottle
##          5              1
##          bow and arrow          box cutter
##          1              13
##          brick              car, knife and mace
##          2              1
##          carjack          chain
##          1              3
##          chain saw          chainsaw
##          2              1
##          chair          claimed to be armed
##          4              1
##          contractor's level    cordless drill
##          1              1
##          crossbow          crowbar
##          9              5
##          fireworks          flagpole
##          1              1
##          flashlight          garden tool
##          2              2
##          glass shard          grenade
##          4              1
##          gun              gun and car
##          3798            12
##          gun and knife    gun and machete
##          22              3
```

##	gun and sword	gun and vehicle
##	1	17
##	guns and explosives	hammer
##	3	18
##	hand torch	hatchet
##	1	14
##	hatchet and gun	ice pick
##	2	1
##	incendiary device	knife
##	2	955
##	knife and vehicle	lawn mower blade
##	1	2
##	machete	machete and gun
##	51	1
##	meat cleaver	metal hand tool
##	6	2
##	metal object	metal pipe
##	5	16
##	metal pole	metal rake
##	4	1
##	metal stick	microphone
##	3	1
##	motorcycle	nail gun
##	1	1
##	oar	pellet gun
##	1	3
##	pen	pepper spray
##	1	2
##	pick-axe	piece of wood
##	4	7
##	pipe	pitchfork
##	7	2
##	pole	pole and knife
##	3	2
##	railroad spikes	rock
##	1	7
##	samurai sword	scissors
##	4	9
##	screwdriver	sharp object
##	16	14
##	shovel	spear
##	7	2
##	stapler	straight edge razor
##	1	5
##	sword	Taser
##	23	34
##	tire iron	toy weapon
##	4	226
##	unarmed	undetermined
##	421	188
##	unknown weapon	vehicle
##	82	213
##	vehicle and gun	vehicle and machete
##	8	1

```
##          walking stick          wasp spray
##                1                1
##          wrench
##                1
```

I am surprised to see air conditioner, flashlight and stapler in the “armed” variable.

Problem 2

- a. Describe the age distribution of the sample. Is this what you would expect to see?

```
hist(dat$age, main= "Age Distribution of Fatal Police Shootings", xlab="Age of Victims", ylab= "Frequency")
```



The distribution of age is skewed to the right. There are some missing values in this distribution. The bulk of the victims were under the age of 40 which makes sense because younger people tend to encounter the police more since they are expected to commit more crimes.

- b. To understand the center of the age distribution, would you use a mean or a median, and why? Find the one you picked.

```
summary(dat$age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##      6.00  27.00   35.00   37.12  45.00   91.00    308
```

Since the age distribution is not a normal bell curve then the median would be a better estimate of the center of the distribution since it is more resistant to skewness. The median of the distribution is 35 years old.

c. Describe the gender distribution of the sample. Do you find this surprising?

```
table(dat$gender)
```

```
##  
##      F      M  
##    3 293 6298
```

```
dat$gender.nas <- ifelse(dat$gender=="", NA, dat$gender)  
table(dat$gender.nas)
```


```
##  
##      F      M  
##   293 6298
```

```
gender.no.nas <- na.omit(dat$gender.nas)  
table(gender.no.nas)
```

```
## gender.no.nas  
##      F      M  
##   293 6298
```

```
tab.gender <- table(gender.no.nas)  
barplot(tab.gender, main= "Gender Distribution of Fatal Police Shootings", xlab= "Gender", ylab="Frequency")
```

Gender Distribution of Fatal Police Shootings



The majority of fatal police shooting victims were males. This is expected because the majority of people targeted by the police are males. Missing values were omitted from the distribution because the gender of those victims were unknown.

Problem 3

- a. How many police officers had a body camera, according to news reports? What proportion is this of all the incidents in the data? Are you surprised that it is so high or low?

```
table(dat$body_camera)
```

```
##  
## False  True  
##  5684   910
```

```
prop.table(table(dat$body_camera))
```

```
##  
##      False      True  
## 0.8619958 0.1380042
```


910 police officers had a body camera according to news reports. This is around 14% of all the incidents in the data. I am surprised that it is so low because I assumed that body cameras have become a norm for police stations but they are expensive and some precincts are not strict with body cameras.

b. In how many of the incidents was the victim fleeing? What proportion is this of the total number of incidents in the data? Is this what you would expect?

```
table(dat$flee)
```

```
##
##           Car      Foot Not fleeing      Other
##      491      1058      845      3952      248
```

```
barplot(table(dat$flee))
```



```
prop.table(table(dat$flee))
```

```
##
##           Car      Foot Not fleeing      Other
## 0.07446163 0.16044889 0.12814680 0.59933273 0.03760995
```

In 1903 incidents the victim was fleeing. This is about 29% of the data. This variable is a bit confusing because of the missing values and the “other” category. It is hard to get an accurate depiction of the data but this is not what I expect because in theory the police would shoot at a person that is a fleeing threat not standing still.


Problem 4 - Answer only one of these (a or b).

- a. Describe the relationship between the variables “body camera” and “flee” using a stacked barplot. What can you conclude from this relationship?

Hint 1: The categories along the x-axis are the options for “flee”, each bar contains information about whether the police officer had a body camera (vertically), and the height along the y-axis shows the frequency of that category).

Hint 2: Also, if you are unsure about the syntax for barplot, run ?barplot in R and see some examples at the bottom of the documentation. This is usually a good way to look up the syntax of R code. You can also Google it.

```
tab.fleecamera <- table(dat$body_camera, dat$flee)
barplot(tab.fleecamera,
        main = "Stacked barchart for Flee and Body Cameras",
        xlab = "Flee", ylab = "Frequency",
        legend.text = rownames(tab.fleecamera),)
```



For all of the categories in “flee”, the majority of police officers did not have a body camera. It appears that the largest proportion of incidents with a body camera are in the “Not fleeing” category however, the “Not fleeing” category has the most observations. Based on this plot it is hard to determine whether or not the victim was actually fleeing because there is no video evidence of the incident.

- b. Describe the relationship between age and race by using a boxplot. What can you conclude from this relationship?

Hint 1: The categories along the x-axis are the race categories and the height along the y-axis is age.

Hint 2: Also, if you are unsure about the syntax for boxplot, run ?boxplot in R and see some examples at the bottom of the documentation. This is usually a good way to look up the syntax of R code. You can also Google it.

Your answer here.

#Extra credit (10 points)

- a. What does this code tell us?

```
mydates <- as.Date(dat$date)
head(mydates)
(mydates[length(mydates)] - mydates[1])
```

This code tells us the dates of all of the incidents and how long the data was collected for by showing the time difference.

- b. On Friday, a new report was published that was described as follows by The Guardian: “More than half of US police killings are mislabelled or not reported, study finds.” Without reading this article now (due to limited time), why do you think police killings might be mislabelled or underreported?

Police killings might be mislabelled or underreported because there is no documentation of the shootings since the majority of officers involved do not have a body camera to record the incident. We have to rely on the police accounts of the incidents which could be biased since the police do not want to show that they were in the wrong. There is also no national system to report all police shootings to which is confirmed by the way that the data was collected for the Washington Post dataset.

- c. Regarding missing values in problem 4, do you see any? If so, do you think that’s all that’s missing from the data?

There is missing data for the variable “Flee”. There are two columns in the “flee” variable that could count under missing data. The first column which is actually missing values and the last column which is “other”. The “other” category is not explained so we do not know what that is reporting. There could potentially be more missing data since the majority of incidents do not have video documentation so the police officers could have lied about the fleeing nature of the victim. The category for body cameras could also have missing values if some police officers did not report whether or not they had a body camera on or they lied.